

- [Home](#)
- [About](#)
- [Contact the Author](#)
- [Follow the Blog](#)
- [License](#)
- [My Projects](#)

Sun 1 Jul 2012

A Simple SQLAlchemy 0.7 / 0.8 Tutorial

Posted by Mike under [Python](#), [SqlAlchemy](#)
[\[3\] Comments](#)

A couple years ago I wrote a rather [flawed tutorial](#) about SQLAlchemy. I decided it was about time for me to re-do that tutorial from scratch and hopefully do a better job of it this time around. Since I'm a music nut, we'll be creating a simple database to store album information. A database isn't a database without some relationships, so we'll create two tables and connect them. Here are a few other things we'll be learning:

- Adding data to each table
- Modifying data
- Deleting data
- Basic queries

But first we need to actually make the database, so that's where we'll begin our journey. Note that SQLAlchemy is a 3rd party package, so you'll need to [install](#) it if you want to follow along.

How to Create a Database

Creating a database with SQLAlchemy is really easy. They have gone completely with their [Declarative](#) method of creating databases now, so we won't be covering the old school method. You can read the code [here](#) and then we'll explain it following the listing. If you want a way to view your SQLite database, I would recommend the SQLite Manager [plugin](#) for Firefox. Or you could use the [simple wxPython application](#) that I created a month ago.

```
# table_def.py
from sqlalchemy import create_engine, ForeignKey
from sqlalchemy import Column, Date, Integer, String
from sqlalchemy.ext.declarative import declarative_base
from sqlalchemy.orm import relationship, backref

engine = create_engine('sqlite:///mymusic.db', echo=True)
Base = declarative_base()

#####
class Artist(Base):
    """
    __tablename__ = "artists"
```

41

Tweet

1

shares

Share

7

reddit

15

```

id = Column(Integer, primary_key=True)
name = Column(String)

#-----
def __init__(self, name):
    """
    self.name = name

#####
class Album(Base):
    """
    __tablename__ = "albums"

    id = Column(Integer, primary_key=True)
    title = Column(String)
    release_date = Column(Date)
    publisher = Column(String)
    media_type = Column(String)

    artist_id = Column(Integer, ForeignKey("artists.id"))
    artist = relationship("Artist", backref=backref("albums", order_by=id))

#-----
def __init__(self, title, release_date, publisher, media_type):
    """
    self.title = title
    self.release_date = release_date
    self.publisher = publisher
    self.media_type = media_type

# create tables
Base.metadata.create_all(engine)

```

If you run this code then you should see the following output sent to stdout:

```

2012-06-27 16:34:24,479 INFO sqlalchemy.engine.base.Engine PRAGMA
table_info("artists")
2012-06-27 16:34:24,479 INFO sqlalchemy.engine.base.Engine ()
2012-06-27 16:34:24,480 INFO sqlalchemy.engine.base.Engine PRAGMA
table_info("albums")
2012-06-27 16:34:24,480 INFO sqlalchemy.engine.base.Engine ()
2012-06-27 16:34:24,480 INFO sqlalchemy.engine.base.Engine
CREATE TABLE artists (
id INTEGER NOT NULL,
name VARCHAR,
PRIMARY KEY (id)
)
2012-06-27 16:34:24,483 INFO sqlalchemy.engine.base.Engine ()
2012-06-27 16:34:24,558 INFO sqlalchemy.engine.base.Engine COMMIT
2012-06-27 16:34:24,559 INFO sqlalchemy.engine.base.Engine
CREATE TABLE albums (
id INTEGER NOT NULL,
title VARCHAR,
release_date DATE,
publisher VARCHAR,
media_type VARCHAR,
artist_id INTEGER,

```

```

PRIMARY KEY (id),
FOREIGN KEY(artist_id) REFERENCES artists (id)
)

2012-06-27 16:34:24,559 INFO sqlalchemy.engine.base.Engine ()
2012-06-27 16:34:24,615 INFO sqlalchemy.engine.base.Engine COMMIT

```

Why did this happen? Because when we created the engine object, we set its **echo** parameter to True. The [engine](#) is where the database connection information is and it has all the DBAPI stuff in it that makes communication with your database possible. You'll note that we're creating a SQLite database. Ever since Python 2.5, SQLite has been supported by the language. If you want to connect to some other database, then you'll need to edit the connection string. Just in case you're confused about what we're talking about, here is the code in question:

```
engine = create_engine('sqlite:///mymusic.db', echo=True)
```

The string, **'sqlite:///mymusic.db'**, is our connection string. Next we create an instance of the declarative base, which is what we'll be basing our table classes on. Next we have two classes, **Artist** and **Album** that define what our database tables will look like. You'll notice that we have Columns, but no column names. SQLAlchemy actually used the variable names as the column names unless you specifically specify one in the Column definition. You'll note that we are using an "id" Integer field as our primary key in both classes. This field will auto-increment. The other columns are pretty self-explanatory until you get to the ForeignKey. Here you'll see that we're tying the **artist_id** to the id in the **Artist** table. The **relationship** directive tells SQLAlchemy to tie the Album class/table to the Artist table. Due to the way we set up the ForeignKey, the relationship directive tells SQLAlchemy that this is a **many-to-one** relationship, which is what we want. Many albums to one artist. You can read more about table relationships [here](#).

The last line of the script will create the tables in the database. If you run this script multiple times, it won't do anything new after the first time as the tables are already created. You could add another table though and then it would create the new one.

How to Insert / Add Data to Your Tables

A database isn't very useful unless it has some data in it. In this section we'll show you how to connect to your database and add some data to the two tables. It's much easier to take a look at some code and then explain it, so let's do that!

```

import datetime
from sqlalchemy import create_engine
from sqlalchemy.orm import sessionmaker
from table_def import Album, Artist

engine = create_engine('sqlite:///mymusic.db', echo=True)

# create a Session
Session = sessionmaker(bind=engine)
session = Session()

# Create an artist
new_artist = Artist("Newsboys")
new_artist.albums = [Album("Read All About It",
                           datetime.date(1988,12,01),
                           "Refuge", "CD")]

```

```
# add more albums
more_albums = [Album("Hell Is for Wimps",
                     datetime.date(1990, 07, 31),
                     "Star Song", "CD"),
                Album("Love Liberty Disco",
                     datetime.date(1999, 11, 16),
                     "Sparrow", "CD"),
                Album("Thrive",
                     datetime.date(2002, 03, 26),
                     "Sparrow", "CD")]
new_artist.albums.extend(more_albums)

# Add the record to the session object
session.add(new_artist)
# commit the record the database
session.commit()

# Add several artists
session.add_all([
    Artist("MXPX"),
    Artist("Kutless"),
    Artist("Thousand Foot Krutch")
])
session.commit()
```

First we need to import our table definitions from the previous script. Then we connect to the database with our engine and create something new, the Session object. The session is our handle to the database and let's us interact with it. We use it to create, modify, and delete records and we also use sessions to query the database. Next we create an Artist object and add an album. You'll note that to add an album, you just create a list of Album objects and set the artist object's "albums" property to that list or you can extend it, as you see in the second part of the example. At the end of the script, we add three additional Artists using the **add_all**. As you have probably noticed by now, you need to use the session object's **commit** method to write the data to the database. Now it's time to turn our attention to modifying the data.

A Note about __init__: As some of my astute readers pointed out, you actually do not need the **__init__** constructors for the table definitions. I left them in there because the official documentation was still using them and I didn't realize I could leave them out. Anyway, if you leave the **__init__** out of your declarative table definition, then you'll need to use keyword arguments when creating a record. For example, you would do the following rather than what was shown in the previous example:

```
new_artist = Artist(name="Newsboys")
new_artist.albums = [Album(title="Read All About It",
                           release_date=datetime.date(1988, 12, 01),
                           publisher="Refuge", media_type="CD")]
```

How to Modify Records with SQLAlchemy

What happens if you saved some bad data. For example, you typed your favorite album's title incorrectly or you got the release date wrong for that fan edition you own? Well you need to learn how to modify that record! This will actually be our jumping off point into learning SQLAlchemy queries as you need to find the record that you need to change and that means you need to write a query for it. Here's some code that shows us the way:

```
from sqlalchemy import create_engine
```

```

from sqlalchemy.orm import sessionmaker
from table_def import Album, Artist

engine = create_engine('sqlite:///mymusic.db', echo=True)

# create a Session
Session = sessionmaker(bind=engine)
session = Session()

# querying for a record in the Artist table
res = session.query(Artist).filter(Artist.name=="Kutless").first()
print res.name

# changing the name
res.name = "Beach Boys"
session.commit()

# editing Album data
artist, album = session.query(Artist, Album).filter(Artist.id==Album.artist_id).filter(
album.title == "Step Up to the Microphone"
session.commit()

```

Our first query goes out and looks up an Artist by name using the **filter** method. The “.first()” tells SQLAlchemy that we only want the first result. We could have used “.all()” if we thought there would be multiple results and we wanted all of them. Anyway, this query returns an Artist object that we can manipulate. As you can see, we changed the **name** from “Kutless” to “Beach Boys” and then committed out changes.

Querying a joined table is a little bit more complicated. This time we wrote a query that queries both our tables. It filters using the Artist id AND the Album title. It returns two objects: an artist and an album. Once we have those, we can easily change the title for the album. Wasn’t that easy? At this point, we should probably note that if we add stuff to the session erroneously, we can rollback our changes/adds/deletes by using **session.rollback()**. Speaking of deleting, let’s tackle that subject!

How to Delete Records in SQLAlchemy

Sometimes you just have to delete a record. Whether it’s because you’re involved in a cover-up or because you don’t want people to know about your love of Britney Spears music, you just have to get rid of the evidence. In this section, we’ll show you how to do just that! Fortunately for us, SQLAlchemy makes deleting records really easy. Just take a look at the following code!

```

# deleting_data.py
from sqlalchemy import create_engine
from sqlalchemy.orm import sessionmaker
from table_def import Album, Artist

engine = create_engine('sqlite:///mymusic.db', echo=True)

# create a Session
Session = sessionmaker(bind=engine)
session = Session()

res = session.query(Artist).filter(Artist.name=="MXPX").first()

session.delete(res)
session.commit()

```

As you can see, all you had to do was create another SQL query to find the record you want to delete and then call `session.delete(res)`. In this case, we deleted our MXPX record. Some people think punk will never die, but they must not know any DBAs! We've already seen queries in action, but let's take a closer look and see if we can learn anything new.

The Basic SQL Queries of SQLAlchemy

SQLAlchemy provides all the queries you'll probably ever need. We'll be spending a little time just looking at a few of the basic ones though, such as a couple simple SELECTs, a JOINed SELECT and using the LIKE query. You'll also learn where to go for information on other types of queries. For now, let's look at some code:

```
# queries.py
from sqlalchemy import create_engine
from sqlalchemy.orm import sessionmaker
from table_def import Album, Artist

engine = create_engine('sqlite:///mymusic.db', echo=True)

# create a Session
Session = sessionmaker(bind=engine)
session = Session()

# how to do a SELECT * (i.e. all)
res = session.query(Artist).all()
for artist in res:
    print artist.name

# how to SELECT the first result
res = session.query(Artist).filter(Artist.name=="Newsboys").first()

# how to sort the results (ORDER_BY)
res = session.query(Album).order_by(Album.title).all()
for album in res:
    print album.title

# how to do a JOINed query
qry = session.query(Artist, Album)
qry = qry.filter(Artist.id==Album.artist_id)
artist, album = qry.filter(Album.title=="Thrive").first()
print

# how to use LIKE in a query
res = session.query(Album).filter(Album.publisher.like("S%a%")).all()
for item in res:
    print item.publisher
```

The first query we run will grab all the artists in the database (a SELECT *) and print out each of their name fields. Next you'll see how to just do a query for a specific artist and return just the first result. The third query shows how to do a SELECT * on the Album table and order the results by album title. The fourth query is the same query (a query on a JOIN) we used in our editing section except that we've broken it down to better fit [PEP8](#) standards regarding line length. Another reason to break down long queries is that they become more readable and easier to fix later on if you messed something up. The last query uses LIKE, which allows us to pattern match or look for something that's "like" a specified string. In this case,

we wanted to find any records that had a Publisher that started with a capital “S”, some character, an “a” and then anything else. So this will match the publishers Sparrow and Star, for example.

SQLAlchemy also supports IN, IS NULL, NOT, AND, OR and all the other filtering keywords that most DBAs use. SQLAlchemy also supports literal SQL, scalars, etc, etc.

Wrapping Up

At this point you should know SQLAlchemy well enough to get started using it confidently. The project also has excellent documentation that you should be able to use to answer just about anything you need to know. If you get stuck, the SQLAlchemy users group / mailing list is very responsive to new users and even the main developers are there to help you figure things out.

Source Code

- [sqlalchemy0708.zip](#)

Further Reading

- SQLAlchemy official [documentation](#)
- wxPython and SQLAlchemy: [Loading Random SQLite Databases for Viewing](#)
- SQLAlchemy: [Connecting to pre-existing databases](#)
- wxPython and SQLAlchemy: [An Intro to MVC and CRUD](#)
- Another Step-by-Step [SQLAlchemy Tutorial](#) (part 1 of 2)



« [Reportlab: Mixing Fixed Content and Flowables](#) | [wxPython by Example Repository](#) »

6 comments · 27 reactions

5 Stars



Leave a message...

Discussion ▾

Community ▾



Florian Kraus · 3 months ago

You've got a little mistake. For getting the first result you have to use the `.first()`-method on a query. Using the `one()`-method will also return only one result, but if there is no entry (or more than one) in the database which matches with your query SQLAlchemy will throw an error. See also

<http://docs.sqlalchemy.org/en/...>

Greetings from Germany :)

1 ^ ▾ · Reply · Share ▸



Beslin · 3 months ago

Nicely explained. Keep posting.

0 ^ | v · Reply · Share ›



Frank · 3 months ago

I m not a technical person so don't know anything about the about post. But may be later I will understand. :)

0 ^ | v · Reply · Share ›



Anon · 3 months ago

Aren't the `__init__` methods in your declarative classes unnecessary? The Base class created by `declarative_base` has a default `__init__` that assigns keyword args to instance attributes. True, you can't pass positional args using the default `__init__`, but in most cases I'm not sure why you'd want to.

I know in your example it doesn't matter much, but IRL I have seen unaware users create massive `__init__` methods just to assign positional args to instance attributes, which was A) a waste of time up front and B) a pain in the ass because you're forced to provide **all** args even if some columns have defaults.

0 ^ | v · Reply · Share ›



Pablo Marti · 3 months ago · parent

They are indeed unnecessary

0 ^ | v · Reply · Share ›



kbracekumar · 3 months ago

There is no need to create `__init__` for sqlalchemy classees, It is added automatically.

<http://docs.sqlalchemy.org/en/...>

0 ^ | v · Reply · Share ›

ALSO ON MOUSE VS. THE PYTHON

What's this? ✕

wxPython: How to Get Children Widgets from aSizer

0 · 1 comment · a month ago



dncarac — Off topic, but I see Cory used the old-fashioned method of calling the `__init__` method of the `wx.Fra...`

Python 101: The Ternary Operator

0 · 8 comments · 25 days ago



Knsqf — This statement will break if condition is true, `on_true` is zero and `on_false` is non-zero: it will as...

Python 101: Exception Handling

0 · 8 comments · 11 days ago



AdamSkutt — If you're going to catch `BaseException`, you might as well do a blanket catch, as I don't really thin...



• Archived Entry

- **Post Date :**
- Sunday, Jul 1st, 2012 at 6:51 pm
- **Category :**
- [Python](#) and [SqlAlchemy](#)
- **Tags :**
- [SQL](#), [SqlAlchemy](#), [sqlite](#)
- **Do More :**
- You can [leave a response](#), or [trackback](#) from your own site.

• Search:

◦

Search

• Links

- - [@MouseVsPython](#)
 - [Python](#)
 - [wxPython](#)

• Archives

◦

Select Month

• Recent Posts

- - [Python 201: List Comprehensions](#)
 - [Tkinter: How to Show / Hide a Window](#)
 - [wxPython: How to Fire Multiple Event Handlers](#)
 - [An Intro to Mercurial](#)
 - [Python PDF Series – An Intro to metaPDF](#)

• Recent Comments

- - !!Dean on [Python 201: List Comprehensions](#)
 - driscollis on [Python PDF Series – An Intro to metaPDF](#)
 - Tim Arnold on [Python PDF Series – An Intro to metaPDF](#)
 - [Mikhail Kashkin](#) on [Python 101: Downloading a File with ftplib](#)
 - [Visto nel Web – 36 « Ok, panico](#) on [Python 101: pip – a replacement for easy install](#)

July 2012

S M T W T F S

[1](#) [2](#) [3](#) [4](#) [5](#) [6](#) [7](#)
[8](#) [9](#) [10](#) [11](#) [12](#) [13](#) [14](#)
 15 [16](#) [17](#) [18](#) [19](#) [20](#) [21](#)
 22 [23](#) [24](#) 25 [26](#) 27 [28](#)
 29 30 31

[« Jun](#)

• Tags

- [binaries](#) [Book Preview](#) [Book Review](#) [Books](#) [ConfigObj](#) [Cross-Platform](#) [databases](#) [Dialogs](#) [Distribution](#) [Django](#) [Education](#) [GUI](#) [IronPython](#) [OOP](#) [Packaging](#) [PSF](#) [PyCon](#) [PyCon 2010](#) [PyCon 2011](#) [PyCon 2012](#) [Pyowa](#) [Python](#) [Python 101](#) [Python Advocacy](#) [Python PDF Series](#) [Python Web Frameworks](#) [PyWin32](#) [Registrv](#) [Reportlab](#) [sizer](#) [Sphinx](#) [SQL](#) [SqlAlchemy](#) [sqlite](#) [System Admin](#) [Testing](#) [TIP](#) [TurboGears](#) [Tutorial](#) [Windows](#) [wxGrid](#) [wxPython](#) [XML](#) [XML Parsing Series](#) [XRC](#)

• Members

- - [Log in](#)
 - [Entries RSS](#)
 - [Comments RSS](#)
 - [WordPress.org](#)

[The Mouse Vs. The Python](#) is powered by [WordPress 3.4.1](#) and delivered to you in 0.895 seconds using 43 queries.
 Theme: [Connections Reloaded](#) by [Ajay D'Souza](#). Derived from [Connections](#).